

Notation

- \mathcal{S} : state space
- \mathcal{A} : action space
- r : reward function
- $\pi(a \mid s)$: policy function
- $p(\tau)$: true (experts') induced distribution over trajectories
- $\tau = ((s_0, a_0), \dots, (s_T, a_T))$: a trajectory of state-action pairs
- $\mathcal{D} = (\tau_1, \dots, \tau_N)$: a dataset of experts' demonstrations

Background

In the MaxEnt framework (Todorov, 2008) of RL, the goal is to find a policy π such that trajectories sampled follow a distribution

$$p(\tau) = \frac{1}{Z} \exp R(\tau)$$

where $R(\tau) := \sum_t r(s_t, a_t)$. The function r is assumed to be deterministic.

Conversely, in inverse RL (within MaxEnt framework) we are presented an experts' policy $\pi^{(e)}$ and want to solve

$$\max_r \mathbb{E}_{\tau \sim p} [R(\tau) - \log Z(r)]$$

Paper 1: GAN-GCL (GAN guided cost learning)

Finn, C., Christiano, P., Abbeel, P., & Levine, S. (2016). A connection between generative adversarial networks, inverse reinforcement learning, and energy-based models. In: NeurIPS 2016.

IRL as MLE

We can interpret it as a maximum likelihood problem

$$\min_{\theta} \mathbb{E}_{\tau \sim p} [-\log p_{\theta}(\tau)]$$

where $p_{\theta}(\tau) = \frac{1}{Z} \exp(-c_{\theta}(\tau))$ is parametrized by the **Boltzmann distribution**.

Typical problem with Boltzmann distribution is estimating the partition function Z . Let $c_\theta(\tau)$ be the energy cost function (think of a negative reward). And assume we can sample from another policy with known density μ . Then

$$Z = \int \exp(-c_\theta(\tau)) d\tau = \int \exp(-c_\theta(\tau)) \frac{\mu(\tau)}{\mu(\tau)} d\tau = \mathbb{E}_{\tau \sim \mu} \left[\frac{\exp(-c_\theta(\tau))}{\mu(\tau)} \right]$$

Then we have a loss function for the parameters θ :

$$\begin{aligned} \mathcal{L}_{\text{IRL}}(\theta) &= \mathbb{E}_{\tau \sim p} [-\log p_\theta(\tau)] \\ &= \mathbb{E}_{\tau \sim p} [c_\theta(\tau)] + \log Z \\ &= \mathbb{E}_{\tau \sim p} [c_\theta(\tau)] + \log \mathbb{E}_{\tau \sim \mu} \left[\frac{\exp(-c_\theta(\tau))}{\mu(\tau)} \right] \end{aligned}$$

The optimal importance sampling for this function is precisely proportional to p [discussion]. We can train any μ to be as close as possible but with a regularization entropy term

$$\mathcal{L}_{\text{sampler}}(\mu) = \mathbb{E}_{\tau \sim \mu}[c_{\theta}(\tau)] + \mathbb{E}_{\tau \sim \mu}[\log q(\tau)]$$

We now take a look at cost-guided GANs and explore the connection.

GANs

For a true data distribution p and generator distribution π the discriminator's objective in GAN is

$$\mathcal{L}_{\text{discriminator}}(D_{\theta}) = \mathbb{E}_{\tau \sim p}[-\log D_{\theta}(\tau)] + \mathbb{E}_{\tau \sim \pi}[-\log(1 - D_{\theta}(\tau))]$$

And the generator objective is

$$\mathcal{L}_{\text{generator}}(\pi) = \mathbb{E}_{\tau \sim \pi}[-\log D_{\theta}(\tau)] + \mathbb{E}_{\tau \sim \pi}[\log(1 - D_{\theta}(\tau))]$$

Cost-guided GANs

For a **fixed** generator with $\pi(\tau)$, the optimal discriminator is

$$D^*(\tau) = \frac{p(\tau)}{p(\tau) + \pi(\tau)}$$

where $p(\tau)$ is the actual distribution of the data.

Moreover, the global minimum of $D_*(\tau)$ is obtained when $\pi \equiv p$, at which the discriminator gives equal probability to fake and generated data.

When the density of π can be evaluated, GAN can be modified to estimate p

$$D_{\theta}(\tau) = \frac{p_{\theta}(\tau)}{p_{\theta}(\tau) + \pi(\tau)} = \frac{\frac{1}{Z} \exp(-c_{\theta}(\tau))}{\frac{1}{Z} \exp(-c_{\theta}(\tau)) + \pi(\tau)}$$

at optimality, $p_{\theta}(\tau) = \frac{1}{Z} \exp(-c_{\theta}(\tau)) = p(\tau)$.

Cost-guided GANs solve IRL

Here is another trick. Let μ be a mixture of the data and policy samples $\mu(\tau) = \frac{1}{2}(p(\tau) + \pi(\tau))$. This μ can be used for the importance sampling estimate:

The authors prove three foundational results:

1. The value of Z which minimizes $\mathcal{L}_{\text{discriminator}}$ is the importance sampling estimate of Z using μ .
2. For this value of Z the derivative, $\partial_{\theta} \mathcal{L}_{\text{discriminator}} = \partial_{\theta} \mathcal{L}_{\text{IRL}}$. Thus the discriminator optimizes \mathcal{L}_{IRL} .
3. The generator's loss satisfies $\mathcal{L}_{\text{generator}}(\pi) = \log Z + \mathcal{L}_{\text{sampler}}(\pi)$.

Conclusion

- Using GAN is equivalent to IRL on trajectories
- Putting a special structure on the discriminator can be used to estimate the expert's policy $p(\tau)$, which in turns gives an estimate of the cost/reward of trajectory
- No examples given by the authors. Why can that be?

Paper 2: Adversarial Inverse Reinforcement Learning (AIRL)

Fu, J., Luo, K., & Levine, S. (2018). *Learning Robust Rewards with Adversarial Inverse Reinforcement Learning*. In: ICLR.

State-action centric vs trajectory-centric

The GAN-GCL of Finn et al. (2016) is trajectory-centric, which gives in high-variance estimates and results in very poor learning.

The goal will be to be able to learn rewards for each state-action pair instead. The problem to solve is **reward entanglement**.

State-action version

Obvious idea is to use GAN-GCL at the state-action level using a discriminator of the form

$$D_{\theta}(s, a) = \frac{\exp(f_{\theta}(s, a))}{\exp(f_{\theta}(s, a)) + \pi(a \mid s)}$$

- f_{θ} serves the experts reward function
- What would happen at optimality?
- Was $\frac{1}{Z}$ necessary, it got out of the picture.

Generator's loss is policy method

Take a look at the generator's loss

$$\begin{aligned}\hat{r}_\theta(s, a) &:= -\log D_\theta(s, a) + \log(1 - D_\theta(s, a)) \\ &= -\log \frac{\exp(f_\theta(s, a))}{\exp(f_\theta(s, a)) + \pi(a \mid s)} + \log \frac{\pi(a \mid s)}{\exp(f_\theta(s, a)) + \pi(a \mid s)} \\ &= -f_\theta(s, a) + \log(\pi(a \mid s))\end{aligned}$$

When adding over trajectories, we obtain the generator's objective

$$\mathcal{L}_{\text{generator}}(\pi) = \mathbb{E}_\pi \left[\sum_{t=0}^T (f_\theta(s_t, a_t) - \log \pi(a_t \mid s_t)) \right]$$

which is an entropy regularized policy method.

Reward ambiguity

- Suppose we learn a reward function $r(s, a, s')$. Then for any $\Phi : \mathcal{S} \rightarrow \mathbb{R}$ a reward function $\hat{r}(s, a, s') = r(s, a, s') + \Phi(s') - \Phi(s)$ leads to the same optimal policy. The term $\Phi(s') - \Phi(s)$ is also called the **shaping term**.
- The actors argue that having a shaped reward function is not good. Because it will not be robust to changes in dynamics. Why?

Entangled rewards

- The reward of a policy should ideally not depend on the environments transition function. But with an entangled reward it does. Why?
- *Example.* Assume deterministic dynamics and $T(s, a)$ the transition function. Then given a different T'

$$r(s, a) + \Phi(T(s, a)) \neq r(s, a) + \Phi(T'(s, a))$$

Theory on entanglement

The authors prove two interesting results

- A reward function is **disentangled** with respect to dynamics if is the same for all transition functions T up to a function of initial state only (I'm skipping technical definition).
- Any reward function $r(s, a, s')$ that is disentangled must be a function of s only.

The algorithmic solution

Since with the proposed methods one cannot learn a function of s only (Why?). The authors proposed the following change to the discriminator

$$D_{\theta,\phi}(s, a, s') = \frac{\exp\{f_{\theta,\phi}(s, a, s')\}}{f_{\theta,\phi}(s, a, s') + \pi(a \mid s)}$$

where

$$f_{\theta,\phi}(s, a, s') = g_{\theta}(s, a) + h_{\phi}(s') - h_{\phi}(s)$$

The shaping term

The **shaping term** h_ϕ has the role of helping mitigate unwanted shaping effects. The authors show that under deterministic dynamics and a ground-truth state-only function

$$h_\phi \mapsto V^*(s)$$

where $V^*(s)$ is the value function. Therefore

$$\begin{aligned} f^*(s, a) &= \{r^*(s) + V^*(T(s, a))\} - V^*(s) \\ &= Q^*(s, a) - V^*(s) \end{aligned}$$

which is the advantage function.

Benchmarks

Table 1: Results on transfer learning tasks. Mean scores (higher is better) are reported over 5 runs. We also include results for TRPO optimizing the ground truth reward, and the performance of a policy learned via GAIL on the training environment.

	State-Only?	Point Mass-Maze	Ant-Disabled
GAN-GCL	No	-40.2	-44.8
GAN-GCL	Yes	-41.8	-43.4
AIRL (ours)	No	-31.2	-41.4
AIRL (ours)	Yes	-8.82	130.3
GAIL, policy transfer	N/A	-29.9	-58.8
TRPO, ground truth	N/A	-8.45	315.5

Paper 3: Variational Adversarial Inverse Reinforcement Learning (VAIRL)

Peng, X. B., Kanazawa, A., Toyer, S., Abbeel, P., & Levine, S. (2019). Variational discriminator bottleneck: Improving imitation learning, inverse rl, and gans by constraining information flow. In: ICLR.

References

- Finn, C., Christiano, P., Abbeel, P., & Levine, S. (2016). A connection between generative adversarial networks, inverse reinforcement learning, and energy-based models. In: NeurIPS.
- Fu, J., Luo, K., & Levine, S. (2018). *Learning Robust Rewards with Adversarial Inverse Reinforcement Learning*. In: ICLR.
- Goodfellow, I., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., & Bengio, Y. (2014). Generative adversarial nets. In: NeurIPS.